LLMs Generate Kitsch

Alexander Koller
Saarland University
koller@coli.uni-saarland.de

Abstract

LLMs are increasingly being used for tasks that have traditionally required human creativity. These include the creation of art, but also activities such as programming and research. I argue that when applied to such tasks, they create kitsch. This is a consequence of being trained first to predict average human artifacts and then to maximize human feedback scores, i.e. to generate inoffensive artifacts of mass appeal.

I would welcome your comments!

1 Introduction

There has been a tremendous amount of recent activity in using large language models (LLMs) and other models of generative AI to produce artifacts that were previously assumed to require creative acts. For instance, GPT 3.5 generates stories and poetry that human annotators rate as highly as human-produced ones (Porter and Machery, 2024); DALL-E generates images that are often indistinguishable from human-generated art and appeals to people; LLMs generate music and even podcasts; and there is an increasing interest in LLMs that can support or even automate the scientific research process (e.g. Jansen et al., 2025).

While the ability to generate computational art is a huge technical advance, these artifacts often leaves behind a feeling of emptiness. The generated music feels formulaic; the generated stories feel "safe" and generic. Generated poetry, such as the example in Figure 1, can recognizably emulate the style of human poets, but often consists of a sequence of shallow tropes. AI-generated art is of increasingly high technical quality (the example poem scans and rhymes well) and often evokes an emotional response in its consumers. But are these responses the kind of novel experience that genuine art provokes, or is the AI-generated work an empty surface that efficiently reflects the human reader's or viewer's own emotions back at them?

She walks the earth with grace and pride,
A beauty that cannot be denied,
With eyes that shine like stars above,
And lips that speak of endless love.

But though she wears a smile so sweet,
A broken heart doth lie beneath,
For in her chest a pain doth beat,
A love unrequited, without relief.

And so she walks with heavy heart,
A figure haunting in the dark,
For love, the sweetest of all art,
Can also leave a painful mark.

Figure 1: An AI-generated poem in the style of Lord Byron, from Porter and Machery (2024).

In this paper, I will argue that content that is generated by today's techniques for generative AI is generally not art, but *kitsch*. Kitsch is widely considered to be bad art, or not even art at all; it describes works that have the superficial characteristics of art, but lack artistic intent, do not enrich the viewer's inner life as real art does, and are designed for mass appeal. While it is difficult to define kitsch precisely, I will review some typical characteristics in Section 2. I will then discuss the ways in which AI-generated art fits these definitions and how our current training regimes promote the generation of kitsch.

My aim in this paper is not to provide a definitive proof of my claim; kitsch is hard to define, and AI-generated art is notoriously hard to evaluate empirically. I instead hope to offer a useful concept for thinking about the differences between art that is produced by humans and by AI, in a time where empirical studies increasingly find it difficult to distinguish between people's responses to them.

Throughout the paper, I will use the word *art* to refer to any activity for which human creators have traditionally required creativity and that are created for the purpose of appreciation by a human *perceiver*. This encompasses text, images, architecture, music, etc. I will broaden my scope to other forms of human creation, such as writing code or doing scientific research, in Section 4.

I am a computational linguist and will therefore mostly talk about LLMs, the type of generative AI with which I am most familiar. However, my arguments are mostly neutral with respect to the modality of the work of art; they apply both to LLMs and to other types of generative AI, and I will include some examples from other modalities.

2 What is kitsch?

When I say "kitsch", any reader of this paper probably has spontaneous examples in mind. Whether we think of garden gnomes or shallow TV shows or overly sentimental music, there is a universal sense among perceivers with even minimal artistic sensibilities that kitsch looks like art on the surface, but something is deeply hollow about it. Kitsch comes across as a deception: it pretends that it is art, but really it evades everything that is deep and meaningful about art and focuses on easy consumption and *mass appeal*.

The word "kitsch" is first attested in the context of the art markets in Munich in the 1870s. This is consistent with the frequent claim that kitsch is a product of the modern era, where the mass appeal of works of kitsch could be monetized through mass production and through distribution in the mass media. Kitsch is generally considered inferior to "real" art, but Ortlieb and Carbon (2019) have argued that the immediate emotional gratification that kitsch offers can give it a positive cognitive value that is simply complementary to that of art.

I will work out some properties of kitsch in more detail, so I can discuss the way in which AI-generated art is kitsch more specifically (Section 3). This section is heavily influenced by Kulka (1996), the most coherent attempt that I could find at actually defining kitsch. All quotes below without an explicit citations are to Kulka's book.

2.1 Appeal to stock emotions

In order to fulfill its purpose of mass appeal, kitsch must address emotions that are shared among a large proportion of the population. It depicts subjects that are capable of arousing such emotions. In the novel *The Unbearable Lightness of Being*, Kundera (1984) describes this connection as follows: "Kitsch causes two tears to flow in quick succession. The first tear says: How nice to see children running on the grass! The second tear says: How nice to be moved, together with all mankind, by children running on the grass! It is the second tear that makes kitsch kitsch."

As a consequence, the emotions that are evoked by kitsch are simple. Children running on the grass are symbolic of the happiness of childhood; the woman in the poem in Figure 1 is sad and in love, and that is all we can take away from it. We can think of Christian kitsch (plastic Jesus babies), nationalist kitsch (fascist speeches), and local kitsch (Eiffel Tower fridge magnets); these evoke a bundle of shallow emotions that many people connect with certain locations, ideologies, and communities they like. A complex, nuanced emotional landscape will not connect as universally with people and therefore damage the work's mass appeal; as Kulka says, kitsch addresses "stock emotions".

The emotional force of kitsch is furthermore derivative of the emotional force of the subject it depicts. People do not buy Eiffel Tower fridge magnets because they find artistic value in them; they buy them to remind themselves of Paris. Similarly, the poem in Figure 1 derives much of its emotional impact from the fact that it is written in the style of Lord Byron, and this evokes the perceiver's prior emotional responses to other works by Lord Byron or other Romantic poets. Kulka argues that the purpose of a work of kitsch is to symbolize the subject it depicts; in his words, "the appeal of kitsch is totally parasitic on the associations related to its referent".

2.2 Competent but conventional surface form

But if the depiction of an emotionally charged subject is the main purpose of a work of kitsch, this means that effective kitsch makes it as easy as possible for the perceiver to identify this subject. Kulka stipulates that kitsch must depict its subject in a way that is "instantly and effortlessly identifiable". This has two important consequences for the surface form of kitsch.

First, a work of kitsch must be executed with a certain degree of *technical competence*. If we attempted to produce kitsch by having a toddler draw a replica of the Mona Lisa, the subject of the work would not be effortlessly identifiable, and

therefore unsuitable as kitsch. Kitschy text must be coherent and grammatically correct; kitschy images must look realistic; kitschy music must involve reasonable chord progressions. Early models of generative AI could not ensure these properties (perhaps image and music generation models still don't). They were instantly recognizable as technically deficient, and therefore nobody seriously took their outputs for either art or kitsch.

Second, "kitsch invariably uses the most *conventional*, standard, well-tried, and tested representational canons." Imagine you wanted to produce kitsch around 1900. You would certainly not invent cubism as a painting style for your work, because this would confuse and perhaps offend your audience, damaging the mass appeal. Instead, kitsch plays it safe; it executes the current conventions competently. Kitsch aims to depict its emotionally charged subject, not to generate artistic innovation.

In other words, just like kitsch is parasitic in terms of *content* on the emotional force of the subject it depicts, it is also parasitic in *style* on other works of art. Other artists must come up with stylistic innovations; once they become the established mainstream, kitsch can then adopt them.

2.3 Kitsch as fake art

"Kitsch is fake art, expressing fake emotions, whose purpose is to deceive the consumer into thinking he feels something deep and serious, when in fact he feels nothing at all" (Scruton, 2014). Similarly, Dorfles (1969) defines kitsch as "something with the external characteristics of art, but which is in fact a falsification of art." There is a sense of deception involved in kitsch: Kitsch pretends to be art, it pretends to enrich our experience with its subject, but it actually leaves us empty.

It seems counterintuitive at first glance that kitsch would be "fake" art, when works of kitsch can in fact be quite beautiful to look at – their purely *aesthetic value* can be high. One can take the position that the value of a work is mainly in its beauty, and all the rest is a matter of taste; this position then accepts that it is difficult to separate art from kitsch.

However, this position misses the point that humans create art for purposes that go beyond beauty. Artists intend to communicate something to the perceiver of the work; art seeks to enrich, enhance, or transform the perceiver's experience with its subject. In addition to its aesthetic value, a work of art has an *artistic value* that has to do with the relevance of this experience and the innovativeness and effectiveness with which it achieves this. By contrast, there is no particular artistic intent behind kitsch, except to evoke the emotions associated with its subject in a way that has mass appeal. We have also seen that kitsch does not innovate with respect to style. Thus, kitsch has very low artistic value. Our experiences with kitsch feel hollow because this low artistic value is executed with a technical competence that makes it look like art on the surface.

3 LLMs generate kitsch

With this clearer understanding of kitsch in mind, I will now argue that when LLMs and other models of generative AI are used to generate art, they generally produce kitsch. This is an almost necessary consequence of the way in which they are currently trained. I assume here that the LLM generates its output without human intervention, and will discuss the human-in-the-loop case in Section 4.

3.1 Conventional surface form

Modern LLMs excel at generating competent surface forms. At least for English, LLMs since GPT-2 have produced reliably grammatical sentences, and since GPT-3, these have mostly formed coherent texts. This competence is the result of pretraining for next-token prediction on very large quantities of data; this allows LLMs to inherit the grammaticality and coherence from the human-written training data. Instruction tuning (Chung et al., 2024; Ouyang et al., 2022) adds the ability to follow instructions specified in a prompt, permitting users to (among other things) specify the style of the surface form.

At the same time, a model that is trained to predict next tokens is primarily a device for generating tokens that are likely in the given context, under the empirical distribution of the training data. This means that LLMs are specifically trained to generate *conventional* surface forms: The training objective rewards the prediction of next tokens that are frequent continuations in the training data. While LLMs are theoretically capable of generating stylistically innovative outputs, these innovative outputs are of low probability, or are even impossible when truncation sampling methods such as top-k sampling are used (Holtzman et al., 2020). Thus, like a work of kitsch, LLM outputs will typically have competent, but conven-

tional surface forms that offer no impediment to recognizing the subject of the generated work.

3.2 Mass appeal

Furthermore, generative AI models are capable of generating outputs with mass appeal. This is illustrated by the current public debate about "AI slop" (Judkis, 2024), cheaply AI-generated art that is being posted en masse to social media and streaming services. While one should distinguish the massive rate of *creation* of AI-generated content from the actual *appeal* of this content to human consumers, there is at least anecdotal evidence for its commercial success. One Youtube channel has 600k subscribers, with individual videos having been viewed 20 million times (Ruwitch, 2025). From a more scientific perspective, studies such as that by Porter and Machery (2024) have found that AI-generated art (in their case, poetry) is rated higher by human subjects than human-written art; it has "appeal" to at least their random sample of the general population.

Again, the mass appeal of LLM-generated art is not an accident – it is directly caused by the training process. In particular, reinforcement learning from human feedback (RLHF, Ouyang et al., 2022) specifically learns a reward model that aims to capture what responses a human user would "like", and then uses it to finetune the LLM to generate highly rated answers. RLHF is an engine for promoting outputs with mass appeal.

3.3 Fake art

As a final point, a work can only be kitsch if it pretends to be art. LLM-generated cooking recipes and travel guides often feel shallow and derivative, due to the conventionality and mass appeal discussed above. But it is only when LLM outputs are presented as art that they evoke the unique sense of morally offensive deception that kitsch does.

There is no doubt that the quantity of overall content that is generated by LLMs has risen sharply since ChatGPT became available. For instance, Liang et al. (2025) estimated that up to 24% of corporate press releases and 10% of job postings in 2024 were written with LLM support. One could argue that corporate text has never claimed to be art, and therefore cannot be kitsch. Indeed, the use of LLMs for these purposes feels like a tragic commentary on the monotony and redundancy of certain white-collar jobs, but it is not offensive from an artistic perspective.

However, an increasing amount of data also suggests a rapid growth of the AI-based production of content that was traditionally considered art. Approximately half of all photos available on Adobe Stock – 313 million as of April 2025 – are self-reported as AI-generated (Kneschke, 2025). The streaming service Deezer (2025a) reported in April 2025 that 18% of all new music uploaded was AI-generated, according to their AI detection tool. This number had risen to 28% in September (Deezer, 2025b). Broadening the scope to other creative activities that are not specifically art, Daniotti et al. (2025) estimate that 30% of Python code uploaded to Github by contributors in the US in December 2024 was AI-generated.

3.4 LLMs generate kitsch

To summarize, generative AI is being used broadly to generate art; its outputs have mass appeal; and the surface form is conventional. This means that AI-generated art satisfies most of the criteria I laid out in Section 2; it is kitsch. This follows from the way in which AI models are being trained.

One criterion which I have not addressed is the reliance of kitsch on the emotional force of the depicted subject and on stock emotions. While this criterion is anecdotally satisfied – the poem in Figure 1 is only one of many examples –, it is difficult to argue for on theoretical grounds. One argument is that LLMs are not humans and therefore cannot have artistic intents that would enrich a human perceiver's inner life; as a consequence, when they choose a subject for the work they generate, they assign a higher probability to subjects that were frequent in the training data.

I hesitate to make this argument too strongly because it feels a bit cheap, and the way in which LLMs recombine and generalize over the contents in the training data is poorly understood. At the same time, if the notion of "stock emotions" could be operationalized, this criterion would lend itself well to empirical evaluation.

4 Kitsch beyond art

Throughout this paper, I have mostly taken a conventional perspective on "art" and therefore "kitsch": human creative activities that produce artifacts for the purpose of appreciation by a perceiver. However, it seems instructive to me to also apply the criteria of kitsch to other forms of human creativity, such as programming and scientific research. These activities are performed to

achieve an external purpose, which is in contrast to art; one can apply notions of "elegance" to code and research, but they are not judged primarily by aesthetic criteria.

Nonetheless, the thoughts in Section 3 apply to the use of LLMs in other creative activities too. We are approaching a point where code generation with LLMs can generate entire apps ("vibe coding") and accelerate human software development, and support both the execution of research and the ideation of new research questions. Following the logic of Section 3, we expect that code and research that is fully generated by LLMs will be conventional and optimized for mass appeal. Kitschy research is commonly called "incremental". Initial empirical studies indicate that purely LLM-based ideation indeed leads to research of lower quality (Si et al., 2025).

This is not to say that generative AI can't play any role in supporting the human creative process. In the hands of a human creator with innovative artistic, scientific, or technical intents, LLMs can both accelerate and inform the creative process. To cite just one prominent recent example, Aaronson (2025) used GPT-5 to provide a key idea to a proof in quantum complexity theory. He recognized a number of GPT-5′s initial suggestions as incorrect, but elicited a useful idea within half an hour of interaction. The literature on using LLMs in research ideation through human-AI collaboration supports this anecdotal finding (Li et al., 2025). It will be interesting to see how future creators will draw the line between those parts of the creative process that requires their own ingenuity and those that they feel comfortable delegating to the machine.

5 Conclusion

In this paper, I have argued that when generative AI is used to replicate human creative activities, it generates kitsch. It is optimized for mass appeal, uses competent but conventional surface forms, and does not enrich the experiencer's inner world with respect to its subject. I have argued that generative AI's propensity to generate kitsch is an almost necessary consequence of its training regime, which combines optimizing the likelihood of the human-produced training data with optimizing appeal to human experiencers.

In the hands of a human with original artistic intentions, generative AI may be able to accelerate and inform their creative process and democratize the production of art. At the same time, online media are being flooded with AI-generated kitsch, which puts human creators at commercial risk.

The outlook that future perceivers of art will be drowned in appealing and technically competent, but parasitic and ultimately soulless kitsch is a dystopia that nobody wants. Are we content to be what Broch (1969) calls a "kitsch-man": a "person who uses this highly considerate mirror so as to be able to recognize himself in the counterfeit image it throws back of him"? It is exactly because generative AI can generate works of such high technical quality that it forces us to engage with the difference between aesthetic and artistic value and the worth of human creation. Perhaps this will lead us to a clearer common understanding of why we, as a society, care about art in the first place.

A note on priority. I first attempted to upload this paper to arXiv on 15 October 2025, days after I discovered that Uhlmann (2025) had also made the connection between LLMs and kitsch. I have worked on the present paper, off and on, since the beginning of September 2024, and I therefore consider both papers to have been developed independently. I have deliberately avoided reading Uhlmann's paper so far in order to prevent accidentally injecting her thoughts into my paper; but I assume it is an insightful complementary take from a humanities perspective and worth reading.

Acknowledgments. I am grateful to Sarah Breckner, Peter Clark, Yupei Du, Sharon Goldwater, Xenia Klinge, and Ellie Pavlick for fruitful and sometimes controversial discussions on the topics in this paper.

References

Scott Aaronson. 2025. The QMA Singularity. Blog post on *Shtetl-Optimized*. Accessed on 15 October 2025.

Hermann Broch. 1969. Notes on the Problem of Kitsch. In Gillo Dorfles, *Kitsch: The World of Bad Taste*. New York: Universe Books.

Hyung Won Chung, Le Hou, Shayne Longpre, Barret Zoph, Yi Tay, William Fedus, Yunxuan Li, Xuezhi Wang, Mostafa Dehghani, Siddhartha Brahma, Albert Webson, Shixiang Shane Gu, Zhuyun Dai, Mirac Suzgun, Xinyun Chen, Aakanksha Chowdhery, Alex Castro-Ros, Marie Pellat, Kevin Robinson, Dasha Valter, Sharan Narang, Gaurav Mishra, Adams Yu, Vincent Zhao,

Yanping Huang, Andrew Dai, Hongkun Yu, Slav Petrov, Ed H. Chi, Jeff Dean, Jacob Devlin, Adam Roberts, Denny Zhou, Quoc V. Le, and Jason Wei. 2024. Scaling Instruction-Finetuned Language Models. *Journal of Machine Learning Research* 25(70):1–53.

Simone Daniotti, Johannes Wachs, Xiangnan Feng, and Frank Neffke. 2025. Who is using AI to code? Global diffusion and impact of generative AI. arXiv: 2506.08945 [cs.CY].

Deezer. 2025a. Deezer reveals 18% of all new music uploaded to streaming is fully AI-generated. Blog post, accessed on 13 October 2025.

Deezer. 2025b. Deezer: 28% of all music delivered to streaming is now fully AI-generated. Blog post, accessed on 13 October 2025.

Gillo Dorfles. 1969. Kitsch. In Gillo Dorfles, *Kitsch: The World of Bad Taste*. New York: Universe Books.

Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. 2020. The Curious Case of Neural Text Degeneration. In *International Conference on Learning Representations*.

Peter Jansen, Oyvind Tafjord, Marissa Radensky, Pao Siangliulue, Tom Hope, Bhavana Dalvi Mishra, Bodhisattwa Prasad Majumder, Daniel S. Weld, and Peter Clark. 2025. CodeScientist: End-to-End Semi-Automated Scientific Discovery with Code-based Experimentation. arXiv: 2503.22708 [cs.AI].

Maura Judkis. 2024. The deluge of bonkers AI art is literally surreal. *Washington Post*, retrieved 25 June 2025.

Robert Kneschke. 2025. Adobe Stock unter Druck: Wie die KI-Bildflut zu neuen Upload-Limits und strengeren Richtlinien führt. Blog post, accessed on 13 October 2025.

Tomas Kulka. 1996. *Kitsch and Art*. Penn State University Press.

Milan Kundera. 1984. *The Unbearable Lightness of Being: A Novel*. HarperCollins.

Sitong Li, Stefano Padilla, Pierre Le Bras, Junyu Dong, and Mike Chantler. 2025. A Review of LLM-Assisted Ideation. arXiv: 2503.00946 [cs.HC].

Weixin Liang, Yaohui Zhang, Mihai Codreanu, Jiayu Wang, Hancheng Cao, and James Zou. 2025. The Widespread Adoption of Large Language Model-Assisted Writing Across Society. arXiv: 2502.09747 [cs.CL].

Stefan A. Ortlieb and Claus-Christian Carbon. 2019. A Functional Model of Kitsch and Art: Linking Aesthetic Appreciation to the Dynamics of Social Motivation. *Frontiers in Psychology* 9.

Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F Christiano, Jan Leike, and Ryan Lowe. 2022. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*, pages 27730–27744.

Bruce Porter and Edouard Machery. 2024. AI-generated poetry is indistinguishable from human-written poetry and is rated more favorably. *Scientific Reports* 14(26133).

John Ruwitch. 2025. 'AI slop' videos may be annoying, but they're racking up views — and ad money. NPR, All Things Considered. Accessed on 15 October 2025.

Roger Scruton. 2014. A Point of View: The strangely enduring power of kitsch. BBC News. Retrieved on 06 June 2025.

Chenglei Si, Tatsunori Hashimoto, and Diyi Yang. 2025. The Ideation-Execution Gap: Execution Outcomes of LLM-Generated versus Human Research Ideas. arXiv: 2506.20803 [cs.CL].

Gyburg Uhlmann. 2025. The Even Sheen of AI: Kitsch, LLMs, and Homogeneity. arXiv: 2509.16794 [cs.CY].